

Nima Kondori

+1 604-396-2858 | nimakondori96@gmail.com | linkedin.com/in/nima-kondori | nimakondori.github.io
Vancouver, BC | Canadian Citizen; TN Eligible; No H-1B Required | Open to Relocation: SF / Seattle / NYC

SUMMARY

Applied Scientist focused on coding-agent evaluation, benchmark design, and LLM-as-judge systems. Designed a multi-axis intent taxonomy, evaluated classifiers, and built an outcome grader that connects offline evaluations with real-world usage signals; combines this work with production ML and multimodal research experience across Python, PyTorch, Azure, and AWS.

EXPERIENCE

Applied Scientist

Apr. 2026 – Present

Microsoft — GitHub

Vancouver, BC

- Designed a multi-axis taxonomy and evaluated LLM classifiers for Copilot agent interactions, achieving 90%+ consistency on each intent axis and enabling offline-vs.-online evaluation gap analysis.
- Adapted intent classification to multi-turn sessions using silver labels from a stronger-LLM panel calibrated with human feedback, then trained a GNN-based model that preserved accuracy against those labels at roughly 2% of the original classifier's cost and latency.
- Analyzed approximately 100k Copilot sessions with the Office of the CTO, identifying distribution gaps between offline evaluations and online usage that informed benchmark-instance selection for agent hillclimbing.
- Designed and built an outcome grader for task completion and user dissatisfaction, adding deeper LLM-based failure analysis for high-dissatisfaction sessions and supporting an in-progress benchmark-generation pipeline.

Senior Machine Learning Scientist

Apr. 2025 – Apr. 2026

Lily AI (Mountain View, CA)

Remote — Vancouver, BC

- Built a multimodal RAG system combining product text, images, and VLM signals, improving attribute-extraction F1 by 25% over legacy methods while prompt and pipeline optimizations reduced token usage by 10% per product.
- Designed SmartQC, an automated quality-control system using 3 LLM judges and MCP-backed tools, reducing human QC cost by 30% through high-confidence attribute auto-approval.
- Fine-tuned transformer-based vision models for image-driven attribute extraction, reducing dependence on text-only pipelines and enabling scalable visual catalog understanding.

Graduate Research Assistant

Sep. 2022 – Apr. 2025

University of British Columbia

Vancouver, BC

- Led the technical work of a 5-person research team developing ControlNet-style video diffusion models for controllable cardiac structure and motion in generated echocardiograms.
- Applied consistency distillation to the video-generation pipeline, reducing latency by 140x from ~7 minutes to ~3 seconds per video while preserving downstream utility.
- Curated a 100k-video echocardiogram dataset and trained ejection-fraction estimation models that achieved 87% accuracy in cardiac-function prediction.

Machine Learning Engineer

Dec. 2022 – Jan. 2024

CanDry Technologies

Vancouver, BC

- Built a ViT-based classifier for production visual quality control, achieving 85% precision and sub-30ms Android inference with TFLite.
- Developed a ControlNet synthetic-data pipeline that generated 1,000 photorealistic minority-class samples and improved rare-state recall by 15%.
- Built a PyTorch Lightning and W&B experimentation platform with 500+ automated hyperparameter trials, halving experiment-iteration time.

Machine Learning Engineer

Oct. 2020 – Sep. 2022

Scenebox

Vancouver, BC

- Built and maintained Python and JavaScript workflows for LiDAR, image, video, geolocation, and time-series data using Airflow, FFmpeg, Kafka, ROS bag processors, Redis, and Elasticsearch.
- Helped rebalance search and identity-management workloads by moving selected Elasticsearch-backed data into SQL storage, reducing query latency by 20% while avoiding unnecessary search-cluster scaling.
- Participated in designing, implementing, and testing a SageMaker Ground Truth integration for enterprise labeling workflows, accelerating annotation turnaround by 30% for 5+ clients.
- Maintained CI/CD infrastructure, deployed Scenebox clusters into customer AWS VPCs, and migrated selected self-managed services to AWS-managed equivalents.

TECHNICAL SKILLS

Agent Evaluation: coding agents, benchmark design, multi-turn agent sessions, telemetry-informed evaluation, LLM-as-judge systems, outcome grading, silver-label calibration

ML and Multimodal: PyTorch, Transformers, Hugging Face, diffusion models, ControlNet, VLMs, multimodal RAG, model distillation, synthetic data

Systems and Data: Airflow, Kafka, Docker, Kubernetes, CI/CD, W&B, MCP, Elasticsearch, PostgreSQL, MySQL

Cloud and Languages: Azure, AWS, Python, C/C++, JavaScript

EDUCATION

University of British Columbia

Vancouver, BC

Master of Applied Science, Electrical & Computer Engineering; GPA: 93.1%

Sep. 2022 – Apr. 2025

- Thesis: Exploring Video Diffusion Models in Echocardiogram Generation

University of British Columbia

Vancouver, BC

Bachelor of Applied Science, Electrical & Computer Engineering; GPA: 89%

Sep. 2017 – May 2020

PUBLICATIONS

HAPPI: Hyperbolic Hierarchical Prototypes for Image Recognition, Beyond Euclidean Workshop, ICCV 2025

ControlEchoSynth: Boosting Ejection Fraction Estimation Models via Controlled Video Diffusion, CVPR 2024 Workshop

ProtoASNet: Dynamic Prototypes for Interpretable and Uncertainty-Aware Aortic Stenosis Classification, MICCAI 2023

SELECTED PROJECT

Agentic Job Finder | *Python, CrewAI, LLM Agents, PostgreSQL*

- Built an agentic resume-tailoring platform with multi-agent validation designed to reduce hallucinated or unsupported claims and YAML-based configuration for reproducible prompt, model, and tool settings.

AWARDS & SERVICE

Awards: Canada Graduate Scholarship (CGS-M), NSERC Undergraduate Student Research Award

Service: Reviewer for IEEE Transactions on Medical Imaging and the NeurIPS 2024 Adaptive Foundation Models Workshop; IEEE TMI Distinguished Reviewer Certificate 2024