Nima Kondori

nimakondori96@gmail.com · linkedin.com/in/nima-kondori · github.com/nimakondori · nimakondori.github.io · Vancouver, BC

Applied Scientist working on evaluation for LLM and agentic systems — intent taxonomies, LLM-as-judge pipelines, benchmark design, and outcome grading — combined with production ML and multimodal research experience across Python, PyTorch, Azure, and AWS.

EXPERIENCE

Applied Scientist Apr. 2026 – Present

Microsoft | GitHub — Vancouver, BC

- Designed a multi-axis taxonomy and evaluated LLM classifiers for Copilot agent interactions, achieving 90%+ consistency on each intent axis and enabling offline-vs.-online evaluation gap analysis.
- Adapted intent classification to multi-turn sessions using silver labels from a stronger-LLM panel calibrated with human feedback, then trained a GNN-based model that preserved accuracy against those labels at roughly 2% of the original classifier's cost and latency.
- Analyzed approximately 100k Copilot sessions with the Office of the CTO, identifying distribution gaps between offline evaluations and online usage that informed benchmark-instance selection for agent hillclimbing.
- Designed and built an outcome grader for task completion and user dissatisfaction, adding deeper LLM-based failure analysis for high-dissatisfaction sessions and supporting an in-progress benchmark-generation pipeline.

Senior Machine Learning Scientist Apr. 2025 – Apr. 2026

Lily AI (Mountain View, CA) — Remote | Vancouver, BC

- Built a multimodal RAG system combining product text, images, and VLM signals, improving attribute-extraction F1 by 25% over legacy methods while prompt and pipeline optimizations reduced token usage by 10% per product.
- Designed SmartQC, an automated quality-control system using 3 LLM judges and MCP-backed tools, reducing human QC cost by 30% through high-confidence attribute auto-approval.
- Fine-tuned transformer-based vision models for image-driven attribute extraction, reducing dependence on text-only pipelines and enabling scalable visual catalog understanding.

Graduate Research Assistant Sep. 2022 – Apr. 2025

University of British Columbia — Vancouver, BC

- Led the technical work of a 5-person research team developing ControlNet-style video diffusion models for controllable cardiac structure and motion in generated echocardiograms.
- Applied consistency distillation to the video-generation pipeline, reducing latency by 140× from ~7 minutes to ~3 seconds per video while preserving downstream utility.
- Curated a 100k-video echocardiogram dataset and trained ejection-fraction estimation models that achieved 87% accuracy in cardiac-function prediction.

Machine Learning Engineer Dec. 2022 – Jan. 2024

CanDry Technologies — Vancouver, BC

- Built a ViT-based classifier for production visual quality control, achieving 85% precision and sub-30ms Android inference with TFLite.
- Developed a ControlNet synthetic-data pipeline that generated 1,000 photorealistic minority-class samples and improved rare-state recall by 15%.
- Built a PyTorch Lightning and W&B experimentation platform with 500+ automated hyperparameter trials, halving experiment-iteration time.

Machine Learning Engineer Oct. 2020 – Sep. 2022

Scenebox — Vancouver, BC

- Built and maintained Python and JavaScript workflows for LiDAR, image, video, geolocation, and time-series data using Airflow, FFmpeg, Kafka, ROS bag processors, Redis, and Elasticsearch.
- Helped rebalance search and identity-management workloads by moving selected Elasticsearch-backed data into SQL storage, reducing query latency by 20% while avoiding unnecessary search-cluster scaling.
- Participated in designing, implementing, and testing a SageMaker Ground Truth integration for enterprise labeling workflows, accelerating annotation turnaround by 30% for 5+ clients.
- Maintained CI/CD infrastructure, deployed Scenebox clusters into customer AWS VPCs, and migrated selected self-managed services to AWS-managed equivalents.

TECHNICAL SKILLS

Agent Evaluation: coding agents, benchmark design, multi-turn agent sessions, telemetry-informed evaluation, LLM-as-judge systems, outcome grading, silver-label calibration

ML and Multimodal: PyTorch, Transformers, Hugging Face, diffusion models, ControlNet, VLMs, multimodal RAG, model distillation, synthetic data

Systems and Data: Airflow, Kafka, Docker, Kubernetes, CI/CD, W&B, MCP, Elasticsearch, PostgreSQL, MySQL

Cloud and Languages: Azure, AWS, Python, C/C++, JavaScript

EDUCATION

Master of Applied Science, Electrical & Computer Engineering Sep. 2022 – Apr. 2025

University of British Columbia — GPA 93.1% · Thesis: Exploring Video Diffusion Models in Echocardiogram Generation

Bachelor of Applied Science, Electrical & Computer Engineering Sep. 2017 – May 2020

University of British Columbia — GPA 89%

PUBLICATIONS

HAPPI: Hyperbolic Hierarchical Part Prototypes for Image Recognition. Beyond Euclidean Workshop, ICCV 2025.

ControlEchoSynth: Boosting Ejection Fraction Estimation Models via Controlled Video Diffusion. CVPR 2024 Workshop.

ProtoASNet: Dynamic Prototypes for Interpretable and Uncertainty-Aware Aortic Stenosis Classification. MICCAI 2023.

SELECTED PROJECT

Agentic Job Finder — Python, CrewAI, LLM Agents, PostgreSQL. Built an agentic resume-tailoring platform with multi-agent validation designed to reduce hallucinated or unsupported claims, and YAML-based configuration for reproducible prompt, model, and tool settings.

AWARDS & SERVICE

Awards: Canada Graduate Scholarship (CGS-M); NSERC Undergraduate Student Research Award.

Service: Reviewer for IEEE Transactions on Medical Imaging and the NeurIPS 2024 Adaptive Foundation Models Workshop; IEEE TMI Distinguished Reviewer Certificate, 2024.