Nima Kondori

linkedin.com/in/nima-kondori | github.com/nimakondori | nimakondori.github.io
Vancouver, Canada | TN Visa Eligible | Open to Seattle / San Francisco / New York

Summary

Applied Scientist focused on coding-agent benchmarks, evaluation systems, and production ML. Hands-on experience improving agent behavior through benchmark analysis, telemetry-informed experiment design, user-outcome grading, multimodal RAG, LLM-as-judge systems, and end-to-end ML infrastructure across Python, PyTorch, Ray, PySpark, Airflow, Kafka, Azure, and AWS.

Experience

Applied Scientist Apr. 2026 – Present
Microsoft / GitHub *Vancouver, BC (Remote)*

- Designed a multi-axis taxonomy and evaluated LLM classifiers for Copilot agent interactions, achieving 90%+ consistency on each intent axis and enabling offline-vs.-online evaluation gap analysis.
- Adapted intent classification to multi-turn sessions using silver labels from a stronger-LLM panel calibrated with human feedback, then trained a GNN-based model that preserved accuracy at roughly 2% of the original classifier's cost and latency.
- Analyzed ∼100k Copilot sessions with the Office of the CTO, identifying distribution gaps between offline evaluations and online usage that informed benchmark-instance selection for agent hillclimbing.
- Designed and built an outcome grader for task completion and user dissatisfaction, adding deeper LLM-based failure analysis for high-dissatisfaction sessions.

Senior Machine Learning Scientist Apr. 2025 – Apr. 2026
Lily AI *Mountain View, CA (Remote)*

- Built a multimodal RAG attribute-extraction system over e-commerce product catalogs, combining text, image, and VLM signals to generate SEO/GEO-optimized product attributes at scale.
- Reduced LLM token usage by 10% per product and improved attribute-extraction F1 by 25% over legacy methods through prompt, data-pipeline, and model-assisted validation optimizations.
- Designed SmartQC, an automated LLM-as-judge quality-control system using 3 independent AI judges and MCP-backed tools; reduced human QC cost by 30% by auto-approving high-confidence attributes.
- Fine-tuned transformer-based vision models for image-driven product attribute extraction, reducing dependency on text-only pipelines and enabling scalable visual catalog understanding.

Graduate Research Assistant Sep. 2022 – Apr. 2025
University of British Columbia *Vancouver, BC*

- Led a 5-person research team building controllable video diffusion models for echocardiogram generation under Dr. Renjie Liao and Dr. Purang Abolmaesumi.
- Trained ControlNet-style video generation models to control cardiac structure and motion dynamics across generated echocardiogram sequences.
- Applied consistency distillation to reduce generation latency by 140x, from ∼7 minutes to ∼3 seconds per video, while preserving downstream utility.
- Curated a 100k-video echocardiogram dataset and trained ejection-fraction estimation models achieving 87% accuracy for downstream cardiac function prediction.

Machine Learning Engineer Dec. 2022 – Jan. 2024
CanDry Technologies *Vancouver, BC*

- Built a ViT-based visual quality-control classifier for production thermal-state detection, achieving 85% precision and sub-30ms Android inference via TFLite.
- Developed a ControlNet-based synthetic data pipeline generating 1,000 photorealistic minority-class samples, improving rare-state recall by 15%.
- Built a PyTorch Lightning + W&B training platform with automated hyperparameter sweeps across 500+ trials, reducing experiment iteration time by 2x.

Machine Learning Engineer Oct. 2020 – Sep. 2022
Scenebox *Vancouver, BC*

- Engineered scalable image and embedding search services using FAISS, achieving sub-300ms retrieval across 10,000+ assets for product and attribute matching workloads.
- Integrated AWS SageMaker Ground Truth annotation pipelines, accelerating annotation turnaround by 30% for 5+ enterprise clients through automated job orchestration.
- Deployed AWS-managed Kafka clusters for real-time ingestion and led an Elasticsearch-to-PostgreSQL migration, reducing message processing time by 10% and query latency by 20%.
- Built Airflow workflows for multimodal data ingestion across image, video, and structured client data, enabling reliable downstream ML training and search pipelines.

Technical Skills

Agentic AI/Evals: coding agents, benchmark design, agent traces, telemetry-informed evaluation, LLM-as-judge systems, user-outcome grading, RAG, tool-use workflows
ML Systems: PyTorch, Transformers, Hugging Face, Diffusers, Ray, PySpark, Airflow, Kafka, Docker, Kubernetes, CI/CD
Agent/Data Stack: LangChain, CrewAI, MCP, Nemotron Agentic Toolkit, FAISS, Qdrant, Elasticsearch, PostgreSQL, MySQL
Cloud/Languages: Azure, AWS, Python, C/C++, JavaScript

Education

University of British Columbia — Vancouver, BC
Master of Applied Science, Electrical & Computer Engineering; GPA: 93.1% — *Sep. 2022 – Apr. 2025*
- Thesis: Exploring Video Diffusion Models in Echocardiogram Generation

University of British Columbia — Vancouver, BC
Bachelor of Applied Science, Electrical & Computer Engineering; GPA: 89% — *Sep. 2017 – May 2020*

Publications

HAPPI: Hyperbolic Hierarchical Prototypes for Image Recognition, Beyond Euclidean Workshop, ICCV 2025
ControlEchoSynth: Boosting Ejection Fraction Estimation Models via Controlled Video Diffusion, CVPR 2024 Workshop
ProtoASNet: Dynamic Prototypes for Interpretable and Uncertainty-Aware Aortic Stenosis Classification, MICCAI 2023

Selected Projects

Browser-Based VLM Agent | *JavaScript, Groq, Llama 3.2-11B, VLMs*
- Built a browser extension that captures user-selected screen regions and routes visual context to VLM agents for structured real-time analysis.
- Implemented multi-agent reasoning workflow with configurable model backends and specialized prompts for visual and textual signals.

Agentic Job Finder | *Python, CrewAI, LLM Agents, PostgreSQL*
- Built an agentic resume-tailoring platform that parses job descriptions and user resumes to generate role-specific resume drafts.
- Designed multi-agent validation workflows to reduce hallucinated or unsupported resume claims and added YAML-based configuration for reproducible prompt/model/tool settings.

Multi-Tasking Transformer | *Python, PyTorch, ViT, Docker*
- Developed a Vision Transformer for cardiac cycle classification and heart-failure estimation, achieving 85% concurrent accuracy across simultaneous tasks.

Awards & Service

Awards: Canada Graduate Scholarship (CGS-M), NSERC Undergraduate Student Research Award, Graduate Support Initiative
Service: Reviewer for IEEE Transactions on Medical Imaging and Adaptive Foundation Models Workshop at NeurIPS 2024; IEEE TMI Distinguished Reviewer Certificate 2024